

Lecture 10

CprE 308

February 2, 2015

Intro

Today's Topics

- Review
- Thread Starting
- Thread Attributes
- Thread Stopping
- Thread Examples

Review

Threads

- Need to have multi-programming
 - Many processes executing in parallel
- Why not use multiple processes?
 - Process creation expensive
 - Each process needs memory, lots of state
 - We don't need all that...
- Multiple “threads” of control within a single process
- Threads share process address space

Process vs. Threads




Figure 1:

The Thread Model

Per process Items

Address space

Global variables

Open files

Child processes

Pending alarms

Signals and signal handlers

Accounting information

Per thread items

Program counter

Registers

Stack

State

Threads

- Threads = Parallelism within the same process
- Will the following benefit from multiple threads?
 - Multiplying huge matrices on a single processor
 - The same problem on multiple processors
 - UNIX Shell

Implementing Threads

Threads Standards

- The POSIX standard describes general thread behavior, and the functions which control threads
 - Individual operating systems are allowed some freedom in how threads are implemented, and hence how they behave
- POSIX 1003.4a -> 1003.1c
 - variants
- Microsoft
 - Win32

Create a pthread

- Use `pthread_create()` to add a new thread of control to the current process.

```
int pthread_create(  
    // The function writes the thread id in tid  
    pthread_t *tid,  
    // The addr. of a thread_attr object is passed in tattr  
    const pthread_attr_t *tattr,  
    // The thread runs a function passed in start_routine  
    void* (* start_routine) (void*),  
    // The start routine's arguments are passed with arg  
    void *arg);
```

Creating a Thread

```
start_servers() {  
    pthread_t thread[nr_of_server_threads];  
    int i;  
    for (i = 0; i < nr_of_server_threads; i++)  
        pthread_create(&thread[i], // thread ID  
                       0, // default attributes  
                       server, // start_routine  
                       argument); // argument  
}
```

```
void *server(void *arg) {  
    while(1) {  
        /* get and handle request */  
    }  
}
```

Example

```
func(int r_in, int r_out, int l_in, int l_out) {  
    pthread_t in_thread, out_thread;  
  
    pthread_create(&in_thread,  
        0,  
        incoming,  
        r_in, l_out);  
    pthread_create(&out_thread,  
        0,  
        outgoing,  
        l_in, r_out);  
}
```

Complications

```
func(int r_in, int r_out, int l_in, int l_out) {  
    pthread_t in_thread, out_thread;  
  
    pthread_create(&in_thread,  
        0,  
        incoming,  
        r_in, l_out); // Can't do this ...  
    pthread_create(&out_thread,  
        0,  
        outgoing,  
        l_in, r_out); // Can't do this ...  
    /* How do we wait till they're done? */  
}
```

Multiple Arguments

```
typedef struct {  
    int first, second;  
} two_ints_t;  
  
func(int r_in, int r_out, int l_in, int l_out) {  
    pthread_t in_thread, out_thread;  
    two_ints_t in={r_in, l_out}, out={l_in, r_out};  
    pthread_create(&in_thread,  
        0,  
        incoming,  
        &in);  
    ...  
}
```

Thread Attributes

Thread attributes

Some POSIX thread (pthread) attributes include:

- A thread may have local or global scope of contention
 - That is, it may compete with all threads in the system for CPU time, or it may compete only with threads in the same task (process)
- A thread has a priority for scheduling
 - Threads may use several scheduling methods, some of which use priority
- A thread may be detached
 - Only non-detached threads may be joined
 - `join` is to wait as thread is to process

The thread attribute object

- The attributes of a thread are held in a thread attribute object, which is a struct defined in `pthread.h`
- You can declare a pthread attribute in your code, but it can only be initialized or modified by the following functions:
 - `int pthread_attr_init(pthread_attr_t *attr);`
 - `pthread_attr_setstackaddr();`
 - `pthread_attr_setstacksize();`
 - `pthread_attr_setdetachstate();`

The thread attribute object (cont)

- Creating a thread using a NULL attribute argument has the same effect as using a default attribute:
 - Non-detached (joinable)
 - With a default stack and stack size
 - With the parent's priority
- To create threads with other attributes, the generating attribute object must be modified using the `pthread_attr_set` functions

Thread Attributes

```
pthread_t thread;  
pthread_attr_t thr_attr;  
  
pthread_attr_init(&thr_attr);  
  
...  
/* establish some attributes */  
...  
pthread_create(&thread, &thr_attr, startroutine, arg);  
...
```

Contrast this approach vs providing a long list of parameters

Stack Size

```
pthread_t thread;  
pthread_attr_t thr_attr;  
  
pthread_attr_init(&thr_attr);  
pthread_attr_setstacksize(&thr_attr, 20*1024*1024);  
...  
pthread_create(&thread, &thr_attr, startroutine, arg);  
...
```

Scheduling policy

```
/* set the scheduling policy to SCHED_OTHER */  
pthread_attr_init(&thr_attr);  
ret = pthread_attr_setschedpolicy(&thr_attr, SCHED_OTHER);  
pthread_create(&thread, &thr_attr, startroutine, arg);
```

Ending threads

Something wrong with this

```
void func(int r_in, int r_out, int l_in, int l_out)
{
    pthread_t in_thread, out_thread;
    two_ints_t in = {r_in, l_out}, out={l_in, r_out};

    pthread_create(&in_thread, 0, incoming, &in);
    pthread_create(&out_thread, 0, outgoing, &out);

    return;
}
```


When Is It Done?

```
void func(int r_in, int r_out, int l_in, int l_out)
{
    pthread_t in_thread, out_thread;
    two_ints_t in = {r_in, l_out}, out={l_in, r_out};

    pthread_create(&in_thread, 0, incoming, &in);
    pthread_create(&out_thread, 0, outgoing, &out);

    pthread_join(in_thread, 0);
    pthread_join(out_thread, 0);
}
```

Waiting for pthreads

- Use `pthread_join()` to wait for a thread to terminate.
- Prototype:

```
int pthread_join(  
    thread_t tid,  
    void **status);
```

- The `pthread_join()` function blocks the calling thread until the thread specified by `tid` terminates. The specified thread must be
 - in the current process, and
 - non-detached

Waiting for pthreads (cont)

- The exit status of the thread specified by `tid` is written to `status` when `pthread_join()` returns successfully.
- Multiple threads cannot wait for the same thread to terminate. If they try to, one thread returns successfully and the others fall with an error of `ESRCH`

Termination

```
pthread_exit((void *) value);  
return((void*) value);  
pthread_join(thread, (void**) &value);
```

Finishing Up

- An important special case arises when the initial thread - the one calling `main()` - returns from `main()` or calls `exit()`.
- This action causes the entire process to terminate, along with all its threads. So take care to ensure that the initial thread does not return from `main()` prematurely.
- Note that when the main thread merely calls `pthread_exit()` it terminates only itself - the other threads in the process, as well as the process, continue to exist
 - The process terminates when all its threads terminates.

Detached Threads

```
start_servers() {  
    pthread_t thread;  
    int t;  
    for ( i=0; i < nr_of_server_threads; i++) {  
        pthread_create(&thread, 0, server, 0);  
        pthread_detach(thread);  
    }  
    ...  
}  
  
server() {  
    ...  
}
```

Examples

Example 1

```
#include <stdio.h>
#include <pthread.h>
#include <string.h>
#define M 3
#define N 4
#define P 5

int A[M][N];
int B[N][P];
int C[M][P];

void *matmult(void *);
```


Example pt 2

```
main() {  
    int i,error; pthread_t thr[M];  
    /* initialize the matrices ... */  
    for( i=0; i<M; i++) { // create the worker threads  
        if (error = pthread_create(  
            &thr[i], 0, matmult, (void *)i)) {  
            fprintf(stderr, "pthread_create: %s", strerror(error));  
            exit(1);  
        }  
    }  
    for( i=0; i<M; i++) // wait for workers to finish jobs  
        pthread_join(thr[i],0);  
    /* print the results ... */  
} // end main
```

Example pt 3

```
void *matmult(void *arg) {  
    int row = (int) arg;  
    int col;  
    int i;  
    int t;  
  
    for(col=0; col < P; col++) {  
        t=0;  
        for( i=0; i<N; i++)  
            t += A[row][i] * B[i][col];  
        C[row][col] = t;  
    }  
    return(0);  
}
```