

Lecture 19 - Deadlocks

CprE 308

February 24, 2013

Review

Review

- Three Examples of Deaclock
 - Multiple Mutex Locks
 - Producer-Consumer Problem
 - Dining PHilosopher Problem

Review: Producer Consumer using Semaphores

Shared Variables

- count (number of items in buffer)
- buffer
- N (maximum size of buffer)

Semaphores

- Empty - semaphore initialized to N (number of free slots in buffer)
- Full - semaphore initialized to zero (number of items in buffer)

Review: Producer Consumer using Semaphores (Example)

Producer

```
while(TRUE) {  
    item = produce();  
    down(Empty);  
    lock(mutex);  
    insert(item,buffer);  
    count++;  
    unlock(mutex);  
    up(Full);  
}
```

Consumer

```
while(TRUE) {  
    down(Full);  
    lock(mutex);  
    item = remove(buffer);  
    count--;  
    unlock(mutex);  
    up(Empty);  
    consume(item);  
}
```

Review: Taking Multiple Locks


Thread A

```
proc1() {  
    pthread_mutex_lock(&m1);  
    /* use object 1 */  
    pthread_mutex_lock(&m2);  
    /* use objects 1 and 2 */  
    pthread_mutex_unlock(&m2);  
    pthread_mutex_unlock(&m1);  
}
```

Thread B

```
proc2() {  
    pthread_mutex_lock(&m2);  
    /* use object 2 */  
    pthread_mutex_lock(&m1);  
    /* use objects 1 and 2 */  
    pthread_mutex_unlock(&m1);  
    pthread_mutex_unlock(&m2);  
}
```

Dining Philosophers



Intro

Deadlocks

Chapter 6

- 1 Resources
- 2 Introduction to Deadlocks
- 3 The Ostrich Algorithm
- 4 Deadlock detection and recovery
- 5 Deadlock avoidance
- 6 Deadlock prevention
- 7 Other issues

Resources

- Examples of computer resources
 - printers
 - tape drives
 - tables
- Processes need access to resources in reasonable order
- Suppose a process holds resource A and requests resource B
 - At same time another process holds B and requests A
 - Both are blocked and remain so

Resources (Cont)

- Deadlocks occur when...
 - Process are granted exclusive access to devices
 - We refer to these devices generally as *resources*
- Preemptable resources
 - Can be taken away from a process with no ill effects
- Nonpreemptable resources
 - will cause the process to fail if taken away

Resource (Cont)

- Sequences of events required to use a resource
 - 1 Request the resource
 - 2 Use the resource
 - 3 Release the resource
- Must wait if request is denied
 - Requesting process may be blocked
 - May fail with error code

Introduction to Deadlocks

- Formal definition:

A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.

- Usually the event is release of a currently held resource
- None of the processes can. . .
 - run
 - release resources
 - be awakened

Four Conditions for Deadlock

1 Mutual Exclusion condition

- each resource assigned to 1 process or is available

2 Hold and wait condition

- process holding resources can request additional

3 No preemption condition

- previously granted resources cannot be forcibly taken away

4 Circular wait condition

- must be a circular chain of 2 or more processes
- each is waiting for resource held by next members of the chain

Modeling

Deadlock Modeling


- Modeled with directed graphs



(a)



(b)



(c)

- resource R assigned to process A
- process B is requesting/waiting for resource S
- process C and D are in deadlock over resources T and U

Deadlock Modeling

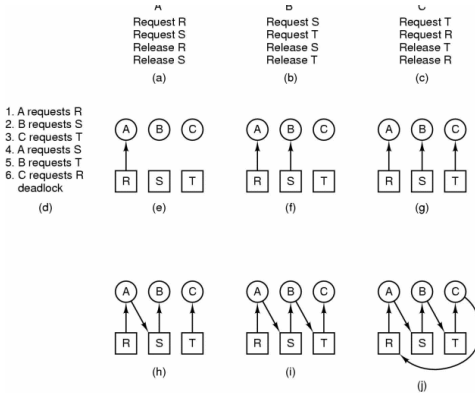


Figure : How deadlock occurs

Deadlock Strategies

Strategies for dealing with deadlocks

- 1 Just ignore the problem altogether
- 2 Detection and recovery
- 3 Dynamic avoidance
 - careful resource allocation
- 4 Prevention
 - Negating one of the four necessary conditions

Ostrich Algorithm

Ostrich Algorithm

Stick your head in the sand

- Pretend there is no problem at all

Detection

Deadlock Detection and Recovery

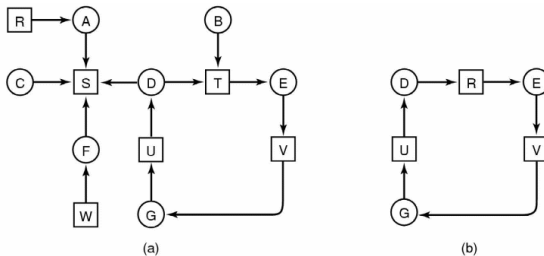
Do not attempt to *prevent* deadlocks

- Instead, detect after the fact and then recover

Deadlock Detection

- Suppose you were given
 - The resource currently held by each
 - The additional resources each process wants
- Question:
 - Is the system deadlocked?

Detection with One Resource of Each Type



- Note the resource ownership and requests
- A cycle can be found within the graph, denoting deadlock

Deadlock Detection Rule (Algorithm)

- Is there some order of execution of the processes which allows all of them to complete?
 - If yes, then not deadlocked
 - If no such order exists, then deadlocked
- After each process finishes, it releases resources back to the system

Example 1

Resource Type	Copies
A	2
B	2
C	2

Process	Resources Held	Resources Needed
P1	A=1, B=1	1 of B
P2	B=1, C=1	1 of C
P3	C=1, A=1	none

Not deadlocked, since P3 can run, releasing resources for P2, which releases resource to P1 upon finishing

Example 2

Resource Type	Copies
A	2
B	2
C	2

Process	Resources Held	Resources Needed
P1	A=1, B=1	1 of B
P2	B=1, C=1	1 of C
P3	C=1, A=1	1 of A

Deadlocked, since nobody can make progress

Example 3

Resource Type	Copies
A	2
B	2
C	2

Process	Resources Held	Resources Needed
P1	A=1, B=1	1 of B
P2	B=1, C=1	2 of A
P3	C=1, A=1	None


Deadlocked:

- Let P3 complete
- P3 releases resources. But then, nobody else can proceed

Detection with Multiple Resources of Each Type

Resources in existence
($E_1, E_2, E_3, \dots, E_m$)

Current allocation matrix




C_{11}	C_{12}	C_{13}	\dots	C_{1m}
C_{21}	C_{22}	C_{23}	\dots	C_{2m}
\vdots	\vdots	\vdots		\vdots
C_{n1}	C_{n2}	C_{n3}	\dots	C_{nm}

Row n is current allocation
to process n

Resources available
($A_1, A_2, A_3, \dots, A_m$)

Request matrix



R_{11}	R_{12}	R_{13}	\dots	R_{1m}
R_{21}	R_{22}	R_{23}	\dots	R_{2m}
\vdots	\vdots	\vdots		\vdots
R_{n1}	R_{n2}	R_{n3}	\dots	R_{nm}

Row 2 is what process 2 needs

Data structures needed by deadlock detection algorithm

Detection with Multiple Resources of Each Type (cont.)

$$E = \begin{pmatrix} 4 & 2 & 3 & 1 \end{pmatrix}$$

Tape drives
Plotters
Scanners
CD Roms

$$A = \begin{pmatrix} 2 & 1 & 0 & 0 \end{pmatrix}$$

Tape drives
Plotters
Scanners
CD Roms

Current allocation matrix

$$C = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 1 \\ 0 & 1 & 2 & 0 \end{bmatrix}$$

Request matrix

$$R = \begin{bmatrix} 2 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 2 & 1 & 0 & 0 \end{bmatrix}$$

An example for the deadlock detection algorithm

Recovery from Deadlock


- Recovery through Preemption
 - Take resource away from owner and give it to another
- Recovery through Rollback
 - Roll back execution to checkpoint
- Recovery through Killing-Process
 - Kill process to break cycle

Avoidance

Deadlock Avoidance

- Be very conservative when granting resources
- Don't grant a resource if it could lead to a *potential deadlock*
- Not very practical, since this is too much overhead for granting resources

Resource Trajectories



Safe and Unsafe States

	Has	Max
A	3	9
B	2	4
C	2	7

Free: 3

(a)

	Has	Max
A	3	9
B	4	4
C	2	7

Free: 1

(b)

	Has	Max
A	3	9
B	0	—
C	2	7

Free: 5

(c)

	Has	Max
A	3	9
B	0	—
C	7	7

Free: 0

(d)

	Has	Max
A	3	9
B	0	—
C	0	—

Free: 7

(e)

Demonstration that the state in (a) is safe

Safe and Unsafe States (cont.)

	Has	Max
A	3	9
B	2	4
C	2	7

Free: 3

(a)

	Has	Max
A	4	9
B	2	4
C	2	7

Free: 2

(b)

	Has	Max
A	4	9
B	4	4
C	2	7

Free: 0

(c)

	Has	Max
A	4	9
B	—	—
C	2	7

Free: 4

(d)

Demonstration that the state in (b) is not safe

The Banker's Algorithm for a Single Resource

Has Max		
A	0	6
B	0	5
C	0	4
D	0	7

Free: 10

(a)

Has Max		
A	1	6
B	1	5
C	2	4
D	4	7

Free: 2

(b)

Has Max		
A	1	6
B	2	5
C	2	4
D	4	7

Free: 1

(c)

- Three resource allocation states
 - Safe
 - Safe
 - Unsafe

Banker's Algorithm for Multiple Resources

	Process	Tape drives	Plotters	Scanners	CD ROMs
A	3	0	1	1	
B	0	1	0	0	
C	1	1	1	0	
D	1	1	0	1	
E	0	0	0	0	

Resources assigned

	Process	Tape drives	Plotters	Scanners	CD ROMs
A	1	1	0	0	
B	0	1	1	2	
C	3	1	0	0	
D	0	0	1	0	
E	2	1	1	0	

Resources still needed

$E = (6342)$
 $P = (5322)$
 $A = (1020)$

Example of a banker's algorithm with multiple resources

Prevention

Deadlock Prevention

- Attacking the mutual exclusion condition
- Attacking the hold and wait condition
- Attacking the no preemption condition
- Attacking the circular wait condition

Attacking the Mutual Exclusion Condition

Somehow allow multiple processes to use resources

- Example: Printer Spooling

Attacking the Hold and Wait Condition

- Require processes to request resources before starting
 - a process never has to wait for what it needs
- Problems
 - may not have required resources at start of run
 - also ties up resources other processes could be using
- Variation:
 - process must give up all resources
 - then request all immediately needed


Attacking the No Preemption Condition

- This is not a viable option
- Consider a process given the printer
 - halfway through its job
 - now forcibly take away printer
 - !????
 - Profit(?)

Attacking the Circular Wait Condition

1. Imagesetter
2. Scanner
3. Plotter
4. Tape drive
5. CD Rom drive

1 2 3



1 2 3

- Normally ordered resources
- A resource graph

Deadlock Prevention summaries

Condition	Approach
Mutual exclusion	Spool everything
Hold and wait	Request all resources initially
No preemption	Take resources away
Circular wait	Order resources numerically

Nonresource Deadlocks

- Possible for two processes to deadlock
 - each is waiting for the other to do some task
- Can happen with semaphores
 - each process required to do a `down()` on two semaphores (*mutex* and another)
 - if done in wrong order, deadlock results

Other Issues

- Two-phase locking
- Communication deadlocks
- Livelock
- Starvation

Communication Deadlocks




Figure : A resource deadlock in a network

Livelock

```
void process_1(void) {  
    enter_region(&resource_1);  
    enter_region(&resource_2);  
    use_both_resources();  
    leave_region(&resource_2);  
    leave_region(&resource_1);  
}
```

```
void process_2(void) {  
    enter_region(&resource_2);  
    enter_region(&resource_1);  
    use_both_resources();  
    leave_region(&resource_1);  
    leave_region(&resource_2);  
}
```