

Lecture 28 - Memory Management

CprE 308

March 24, 2014

Memory Management

Overview

Ideal World (for the programmer):

- I'm the only process in the world
- I have a huge amount of memory at my disposal

Real World

- Many processes in the system
- Not enough memory for them all

Goal of Memory Management

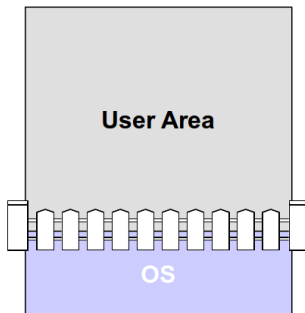
- Present the ideal world view to the programmer, yet implement it on a real system

Simple Memory Management

- One program at a time (not ideal for performance)
- What does running a program involve?
 - Load program into memory
 - Jump to the first instruction
- Issues:
 - Need to protect the OS from the user program
 - Relocation: User memory references should lead to real memory locations

Memory Fence

Memory Fence



The Relocation Problem


- Your Program is complied and linked
 - Generates (absolute) addresses
- When being loaded into memory, can't predict which address it will be loaded in
- How to ensure correct memory addresses are used, and to manage protection?
- One solution: at loading time, update all the addresses to reflect correct physical addresses
 - Still does not ensure protection

Address Space


The Address Space Concept

- Protect processes from one another
- Protect the OS from user processes
- Provide efficient management of available storage

Base and Bounds Registers



Swapping



Address Space

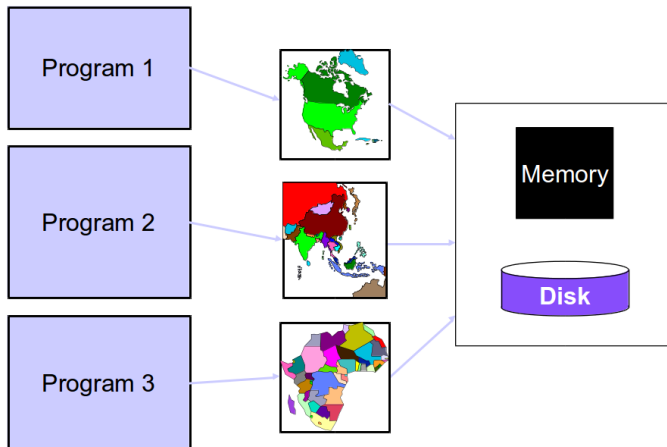
- Each process has it's own “address space” - set of addresses the program generates
- Memory Manager maps the address space onto Physical Memory + Secondary Storage
- Address spaces of different processes don't overlap

Virtual Memory

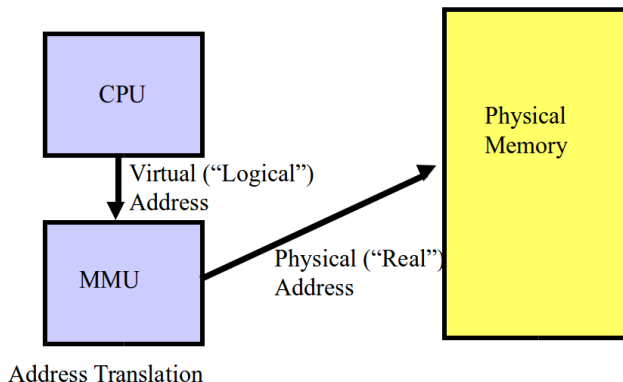
Advantages of Virtual Memory

- Program size can be larger than physical memory
 - Some pages in memory, some on the disk
- Processes can run fast as long as all their active pages are in memory (compare with swapping)

Virtual Memory



Memory Management Unit (MMU)



Virtual memory

- 1 Process generates a “virtual address”
- 2 Virtual address is translated into a physical address
(translation is more sophisticated than in base + bound registers)
 - Base + Bound: add and compare for every memory reference
(Translation done differently for different processes)
- 3 Physical address goes onto the bus

Structuring Virtual Memory

- Paging
 - divide the address space into fixed-size pages
 - Internal fragmentation
- Segmentation
 - divide the address space into variable-size segments (typically each corresponding to some logical unit of the program, such as a module or subroutine)
 - External fragmentation
- Paged segmentation

MMU - a closer look

Lookup the virtual address

- If present in physical memory, they fetch it
- If not, then call upon the operating system (“page fault”)
 - OS loads the required page into memory from secondary storage
 - Instruction is re-started

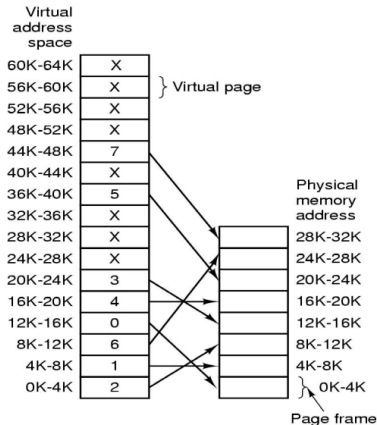
Address Translation

Address Translation Problem

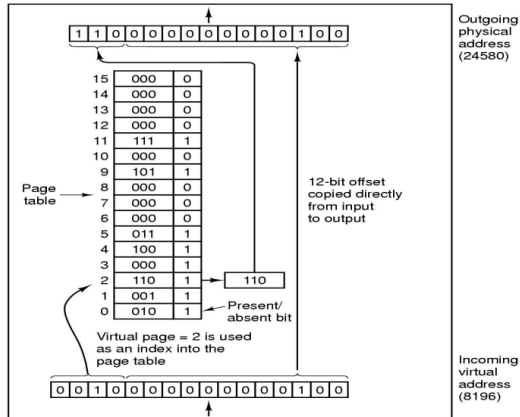
- Given the virtual address
- To find:
 - If in memory, the physical address
 - Else, answer saying “not in memory”
- Constraint: Space overhead of translation should be small

Page Tables

- 64K Virtual Address Space
 - Divided into 16 pages, of 4K each
- 32K Physical Memory
 - 8 pages of 4K each



Page Table




Page Fault

What if required page not in memory?

- 1 Page fault gives control to the OS
- 2 OS fetches it from the disk
 - Needs to evict an existing page from memory (page replacement policy)
- 3 Instruction is restarted

Typical Page Table Entry



Performance

- The address translation is done on every memory reference
- Maybe twice per instruction
 - Instruction fetch
 - Fetch Memory operand
- The translation better be fast!

Different Schemes for Address Translation

- Page Table - most popular
- Translation Lookaside Buffer (TLB)
 - Performance reasons (cache)
- Inverted Page Table
 - Large address spaces

Memory Management with Linked Lists

- Maintaining a linked list of allocated and free memory segments
- First fit
- Best fit
- Worst fit

Question

Where is the page table stored?

- Registers?
- Memory?