

Lecture 37 - File Systems: Shared Files, Performance

CprE 308

April 14, 2014

Intro

Today's topics: File System Implementation

- Shared Files
- Buffer Cache and File System Consistency


Sharing

Sharing of Files


- In UNIX: `ln src dest`
- Two ways of linking files
 - “hard” links
 - Symbolic links

Hard links

- Both files point to the same inode
- `ln /home/guan/f1 /home/guan/f2`



Hard Links




(a) Situation prior to linking

(b) After the link is created

(c) After the original owner removes the file

Symbolic links

- Files point to different inodes
- `ln -s /home/guan/f1 /home/guan/f2`






Performance

Performance of File System

- Where does your data go after a `write()` system call?
- Where does the data come from for a `read()`?
- Think about performance

The Buffer Cache




Buffer Cache

Read(block)

- 1 See if block present in buffer cache
 - If yes, then return buffer
- 2 Initiate disk read for the block
- 3 Sleep till read is complete
- 4 Return buffer

Read Ahead



Buffer Cache - Write

`Write(block) // assume block in cache`

- (Usually) Write to cache and return; the write to disk is done later ([write-back cache](#))
- (Sometimes) Write to cache, schedule a write to disk and return ([write-through cache](#))
- (Exceptional cases) Write to cache, do a synchronous (blocking) write to disk, and return


Write

- Write-back more efficient than write-through
- A disk crash might cause a more serious problem with write-back
- What happens when:
 - The system is turned off without a shutdown
 - A floppy is removed from the drive without unmounting
- System to the rescue: Every 30 seconds or so, a `sync` is done, writing all cache contents to disk


Structure of Cache

- Memory allocated by the system
- Lookup:
 - hash tables
- Page Replacement:
 - LRU
 - Keep a list sorted according to time of use

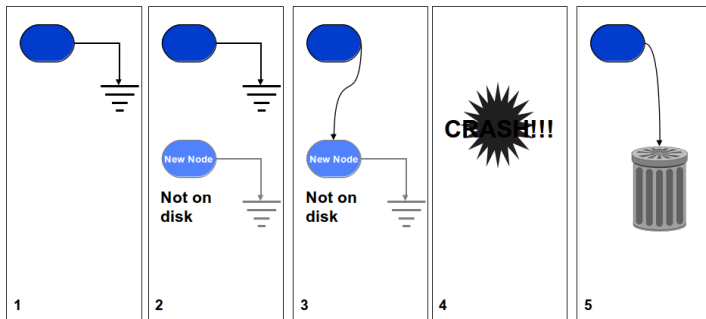
Structure of the Cache




File-System Consistency (1)



File-System Consistency (2)



Keeping It Consistent



File Systems do Crash

- Bring to a consistent state using `fsck` on Unix
- Make sure every disk block is in exactly one file or on the free list
- Go through all directories, and count the number of links per file - check for inconsistencies
- Might prompt the user before taking action

Log Structured File Systems

- If there's lots of caching, then most operations to the file system are *writes*
- Writes are quickest when there is no need to do a seek
- Thus: perform writes wherever the disk head happens to be

